JetPack, Scikit-learn Installation

RTLab Seongsu Keum

CONTENTS

1 JetPack Installation
- host install
- target install
- car detect sample

2 Scikit-learn Installation
- numpy, scipy, matplotlib
- scikit learn

3 NVIDIA Two Days to a Demo
- ...

01 JetPack Installation

JetPack

NVIDIA JetPack SDK is the most comprehensive solution for building AI applications. Use the JetPack installer to flash your Jetson Developer Kit with the latest OS image, install developer tools for both host PC and Developer Kit, and install the libraries and APIs, samples, and documentation needed to jumpstart your development environment.

Key Features in JetPack

- TensorRT
- cuDNN
- CUDA
- VisionWorks/OpenCV
- Multimedia API
- L4T
- Developer Tools

01 JetPack Installation

- host PC (Ubuntu 16.04)
- NVIDIA Jetson TX2
- Internet accessible router/switch

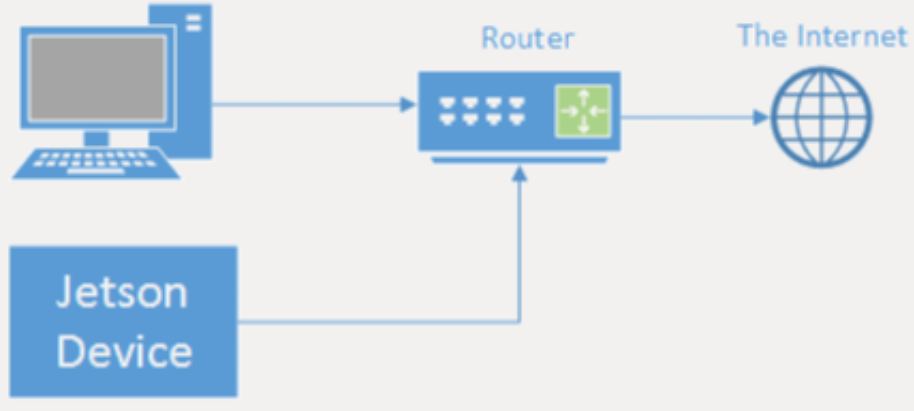

Jetpack (v3.0) Install Guide
https://www.youtube.com/watch?v=D7lkth34rgM

01 JetPack Installation

https://developer.nvidia.com/embedded/jetpack

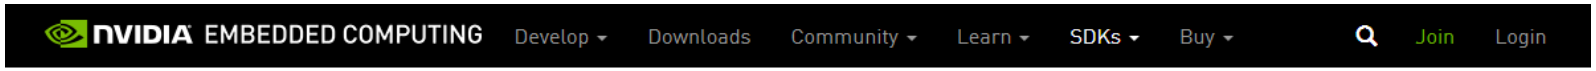

Home > Autonomous Machines > Develop > Tools > JetPack

JetPack

NVIDIA JetPack SDK is the most comprehensive solution for building AI applications. Use the JetPack installer to flash your Jetson Developer Kit with the latest OS image, install developer tools for both host PC and Developer Kit, and install the libraries and APIs, samples, and documentation needed to jumpstart your development environment.

JetPack 4.1 Developer Preview *Early Access*

JetPack 4.1 Developer Preview Early Access is an update to our early access release. It contains a security update and includes functional and performance enhancements for NVIDIA Jetson AGX Xavier Developer Kit, including:

- Improved bandwidth utilization of the memory subsystem
- Enables coherency for Ethernet transactions
- Fixes the installer to not fail silently when installing from unsupported file systems

See Highlights below for a summary about new features enabled with this release, and view the JetPack release notes for more details, including information about additional functionality planned for future releases.

Download JetPack 4.1 Developer Preview EA, and view the full release notes here.

Read the Jetson AGX Xavier Developer Kit User Guide here.

JetPack 3.3

JetPack 3.3 remains the current production release, supporting Jetson TX2 and Jetson TX1 Developer Kits.

Download JetPack 3.3, and view the full release notes here.

01 JetPack Installation

Download JetPack

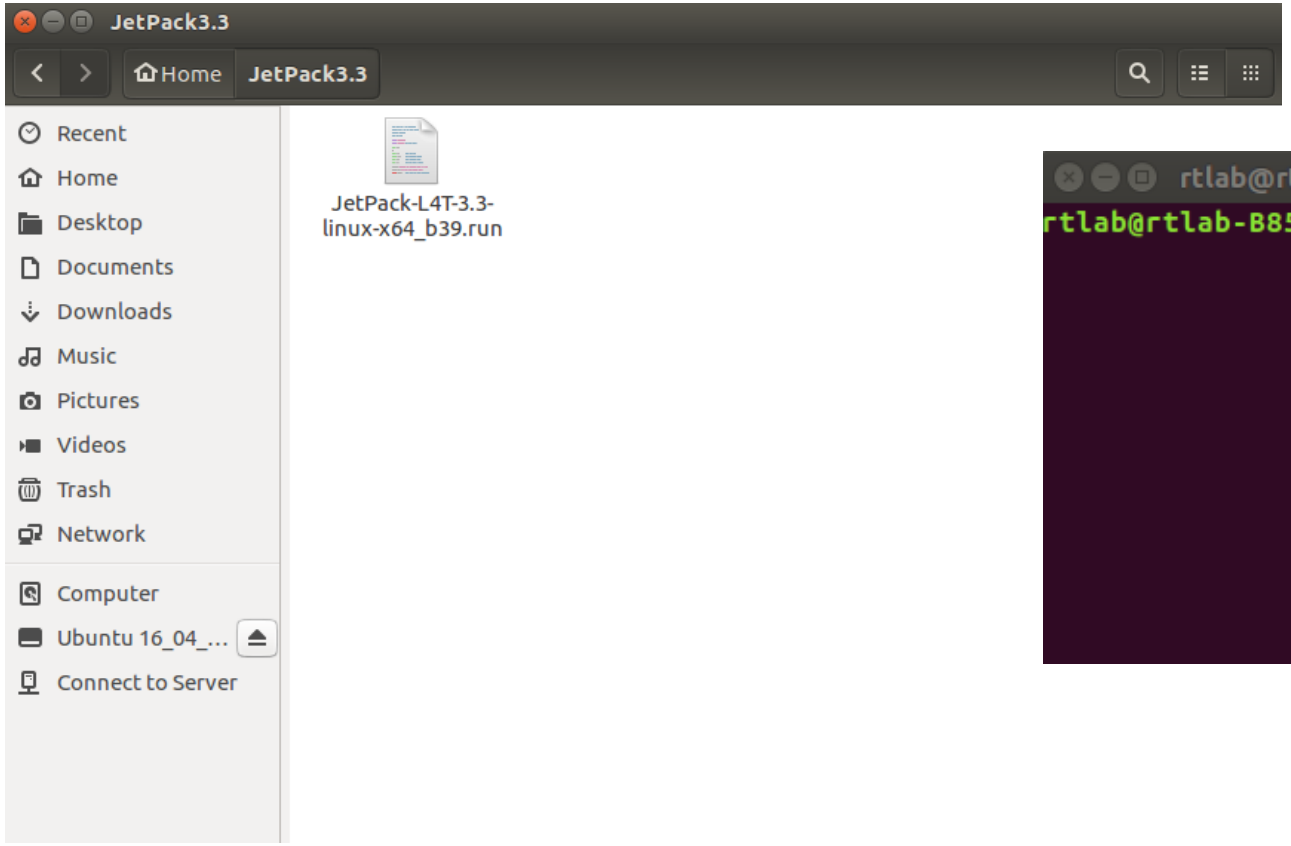

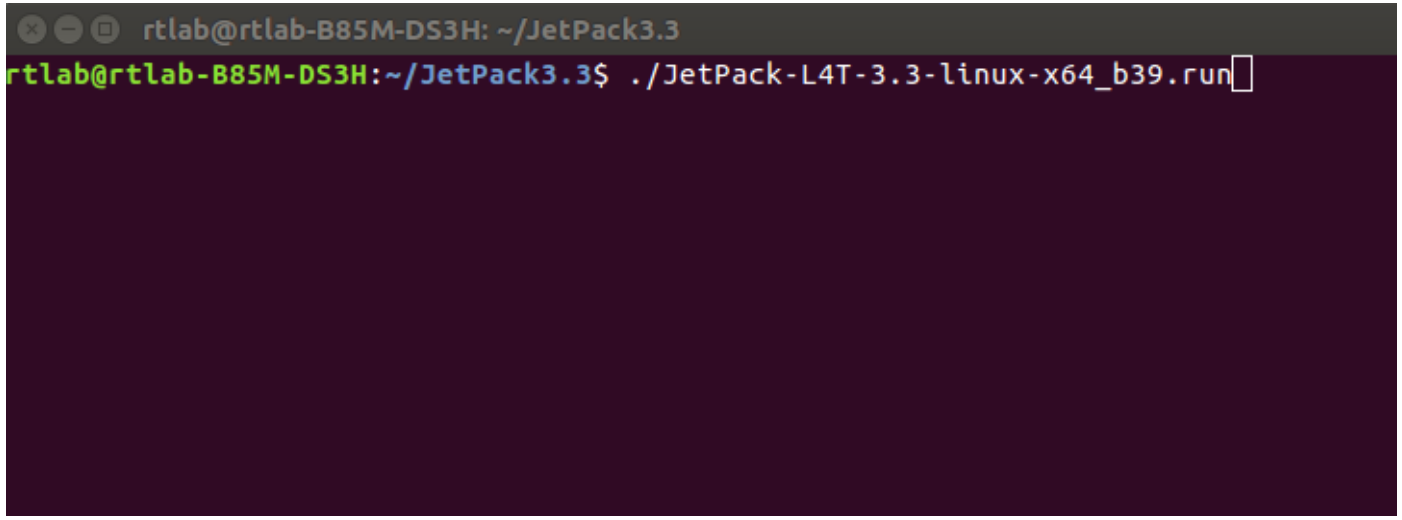

01 JetPack Installation

Jetpack Installer

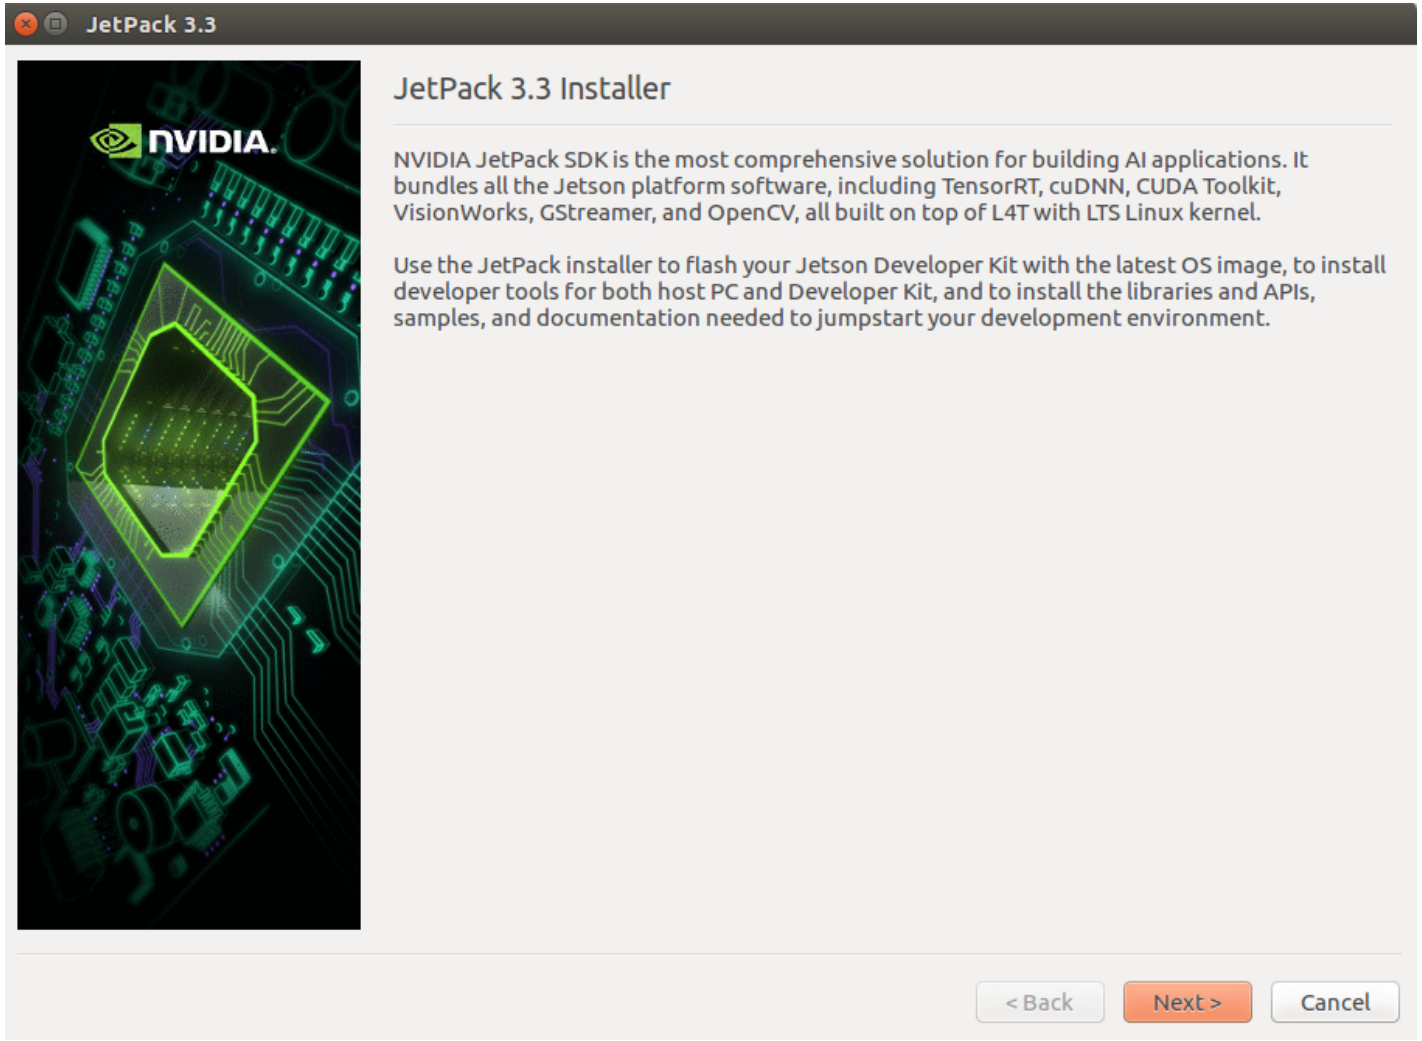

01 JetPack Installation

Install Packages

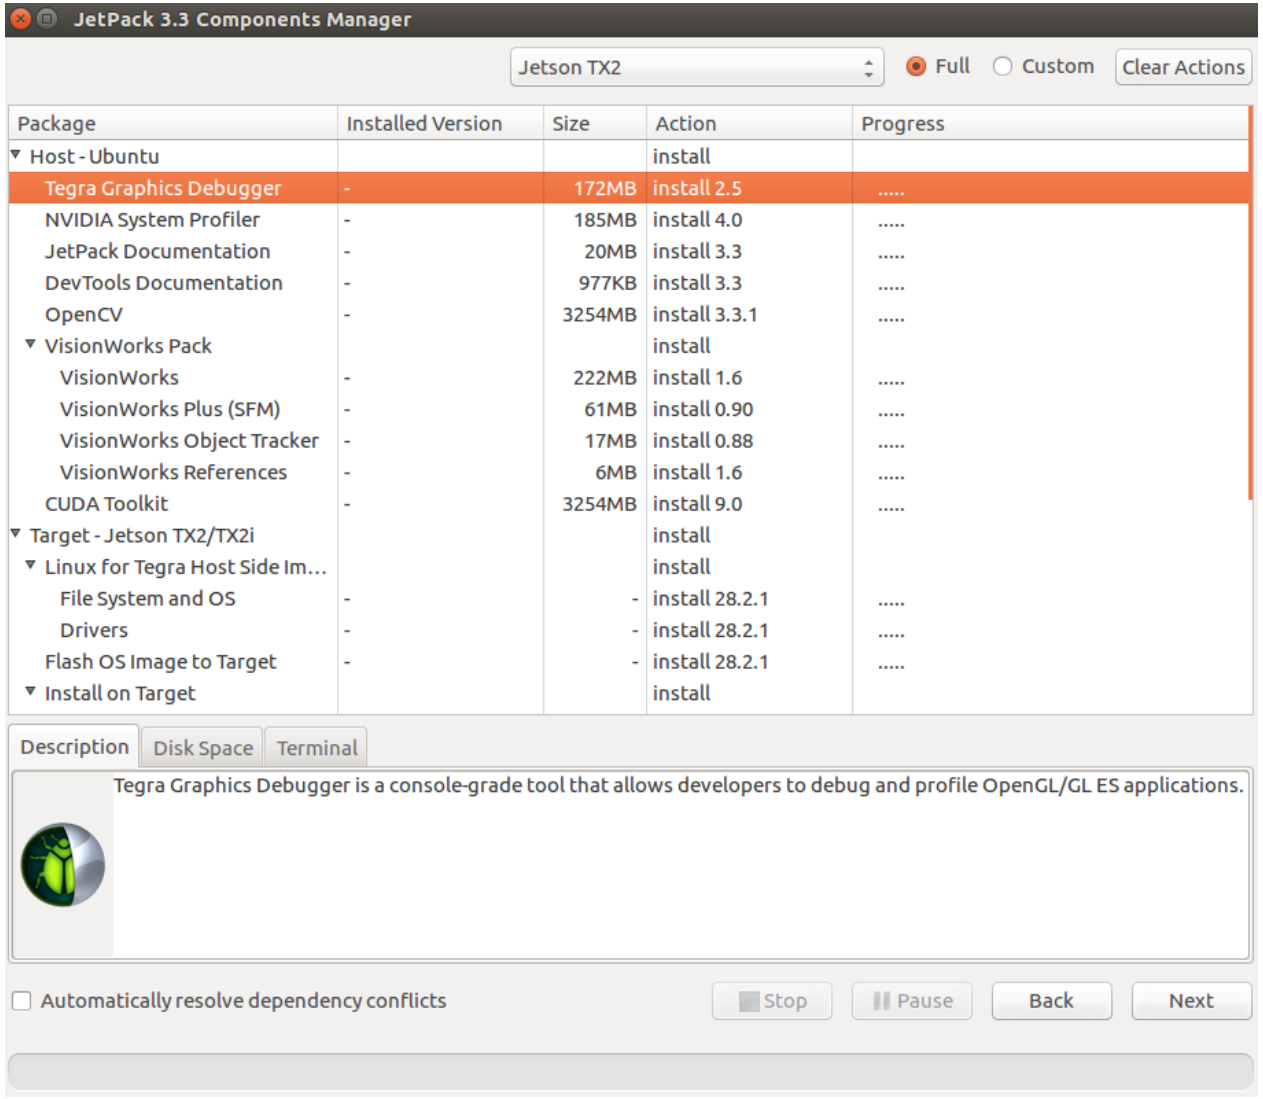

Accept License

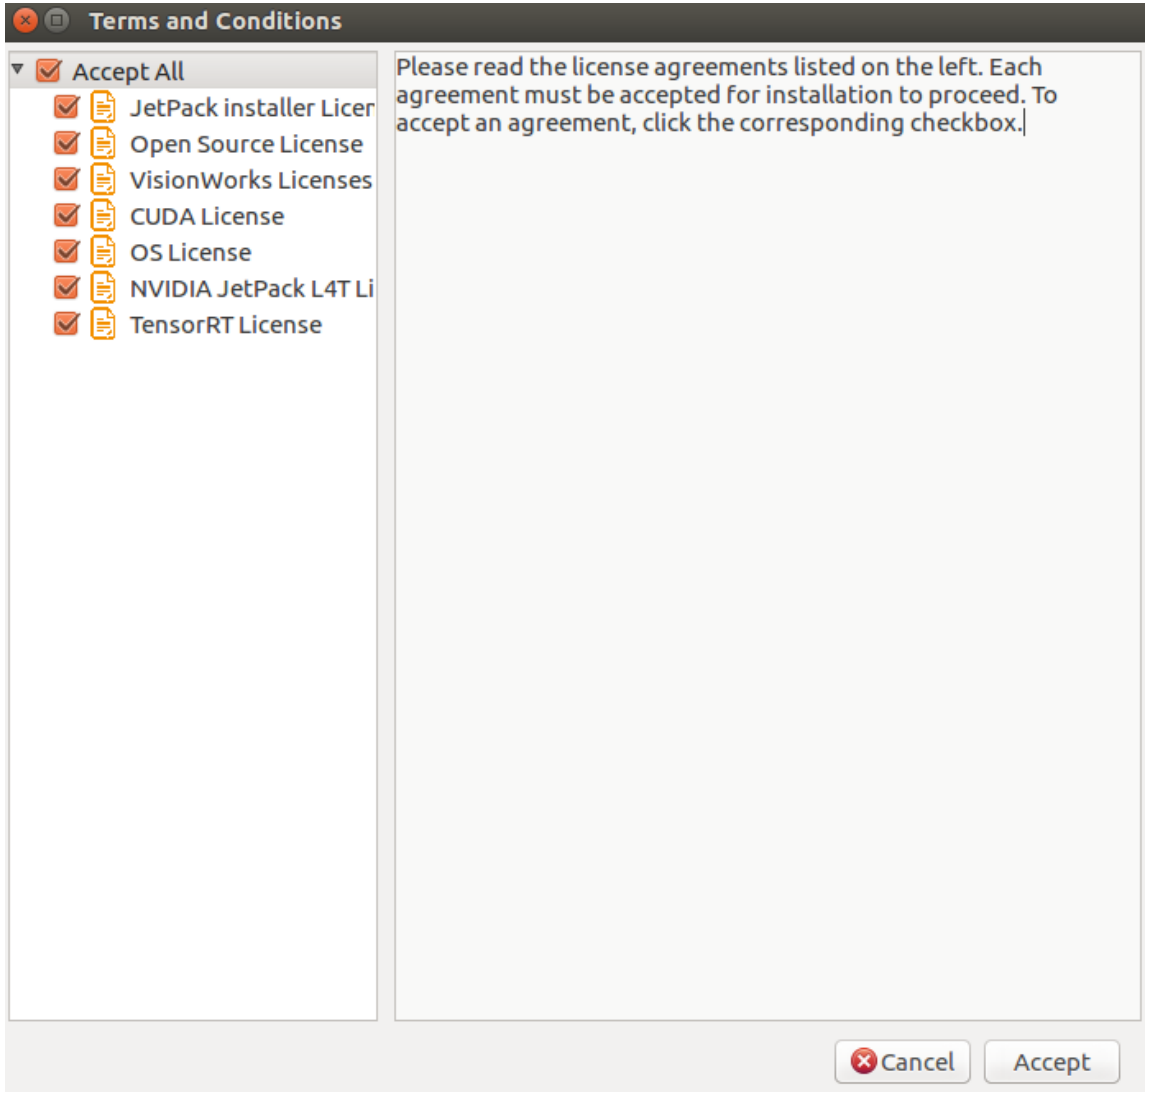

01 JetPack Installation

Host Installation took ~15 minutes

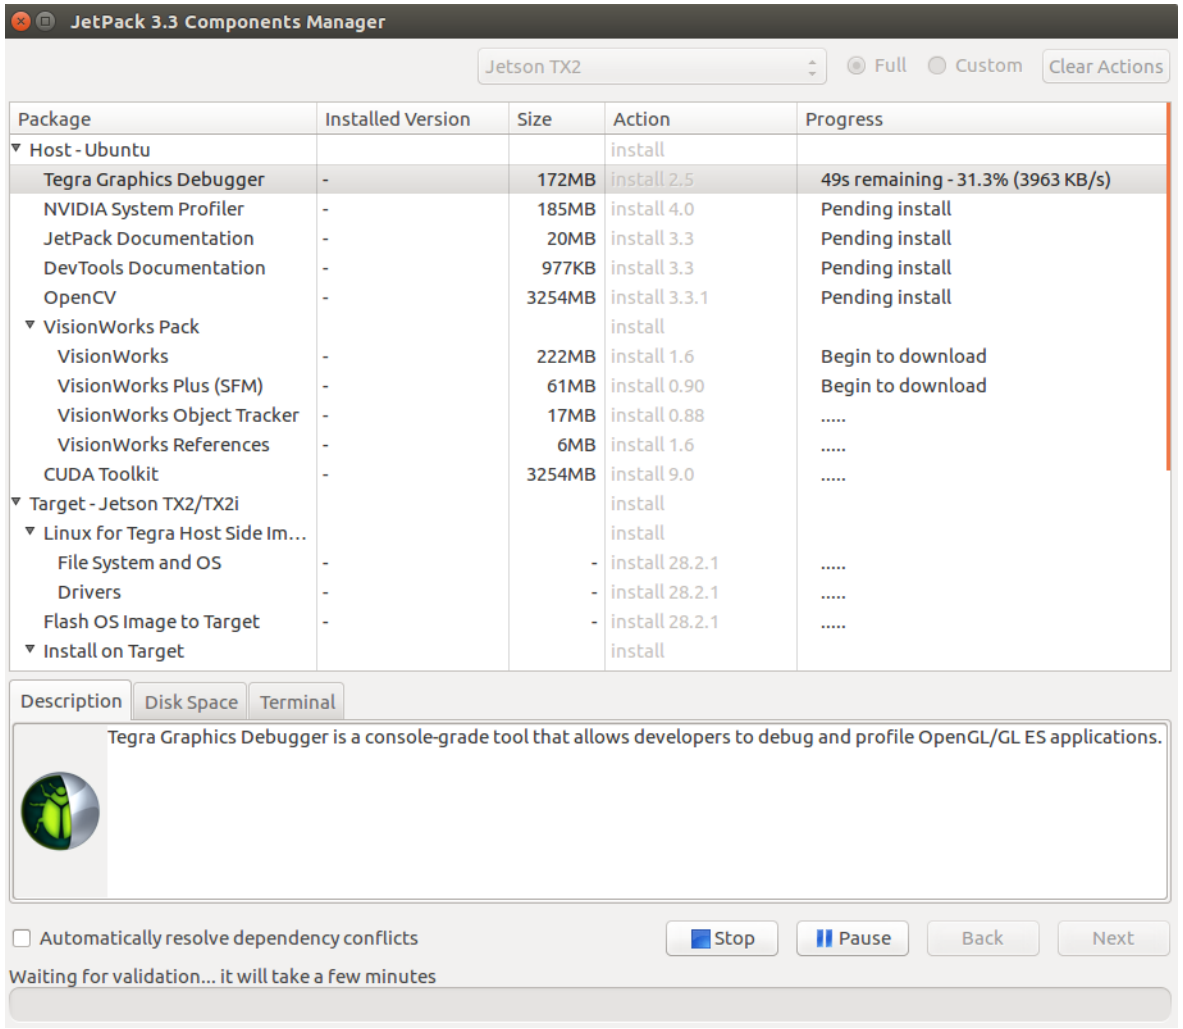

Completion of Host Installation

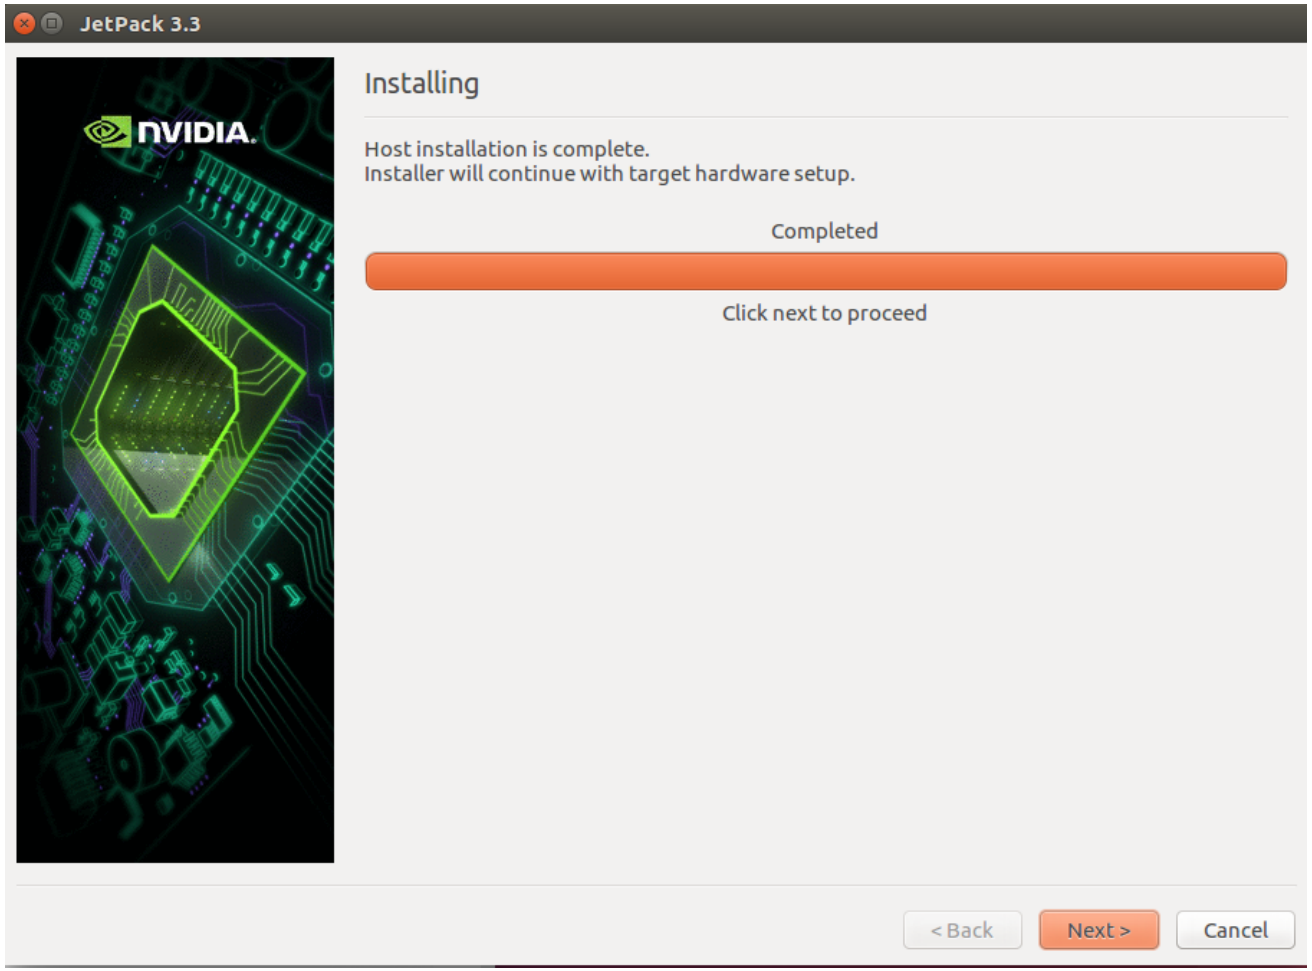

01 JetPack Installation

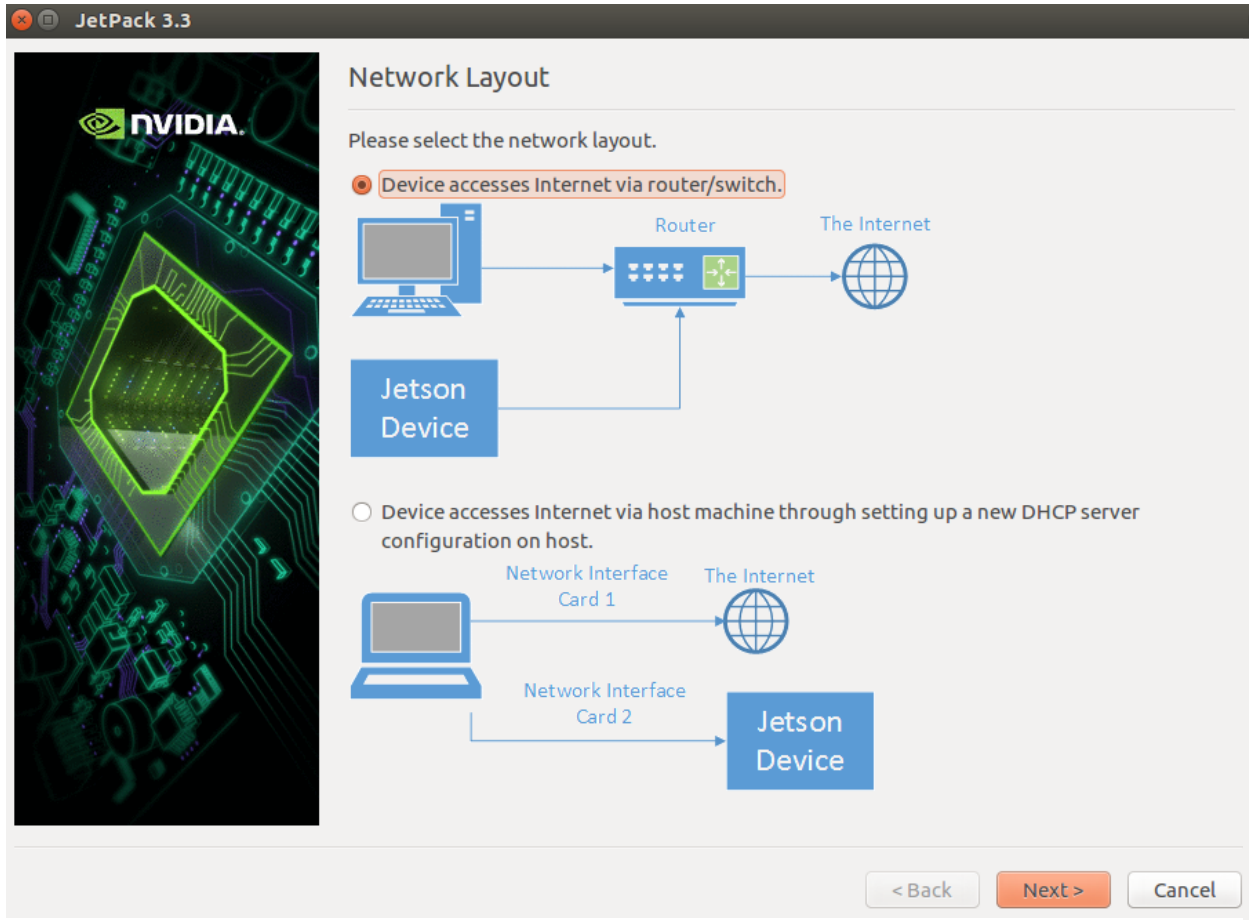

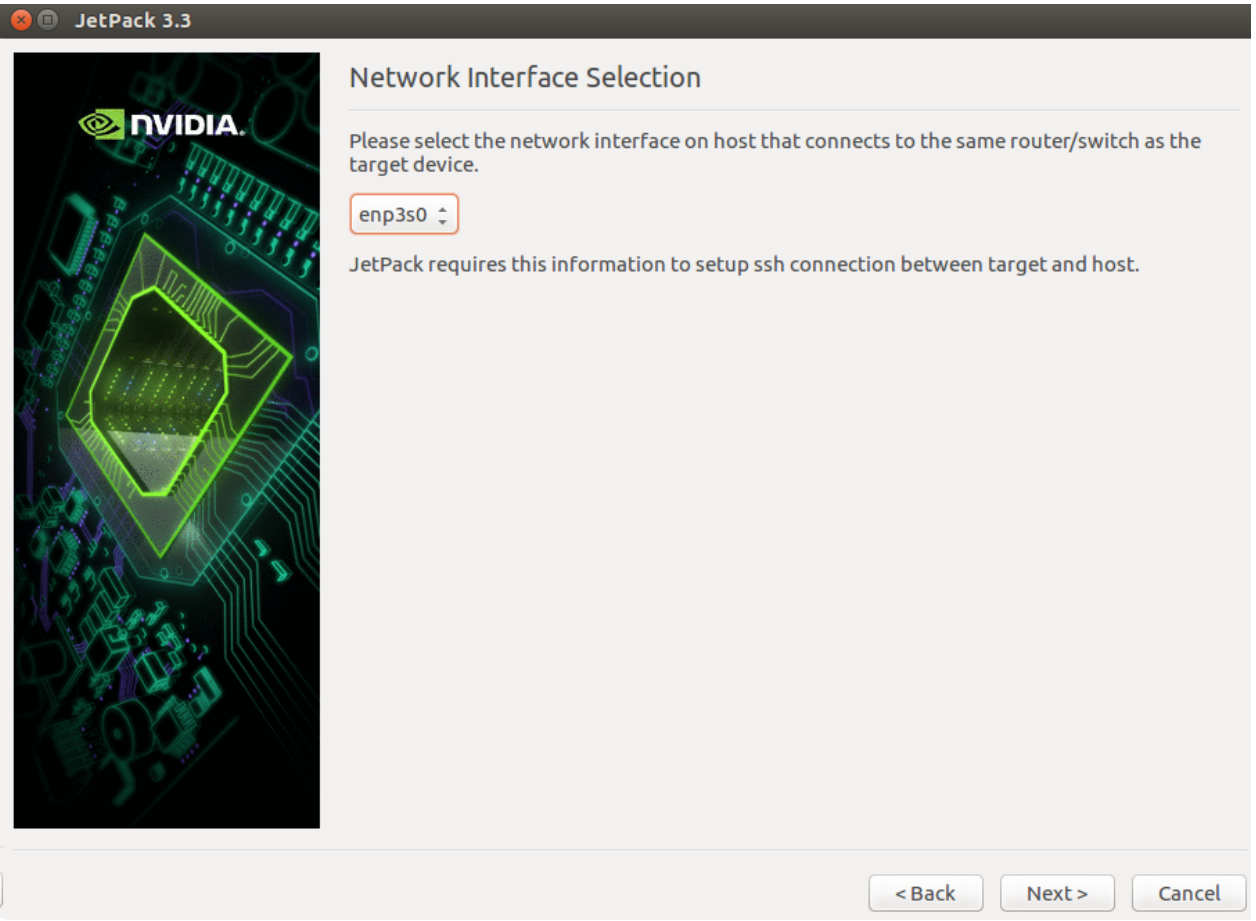

01 JetPack Installation

Target Installation

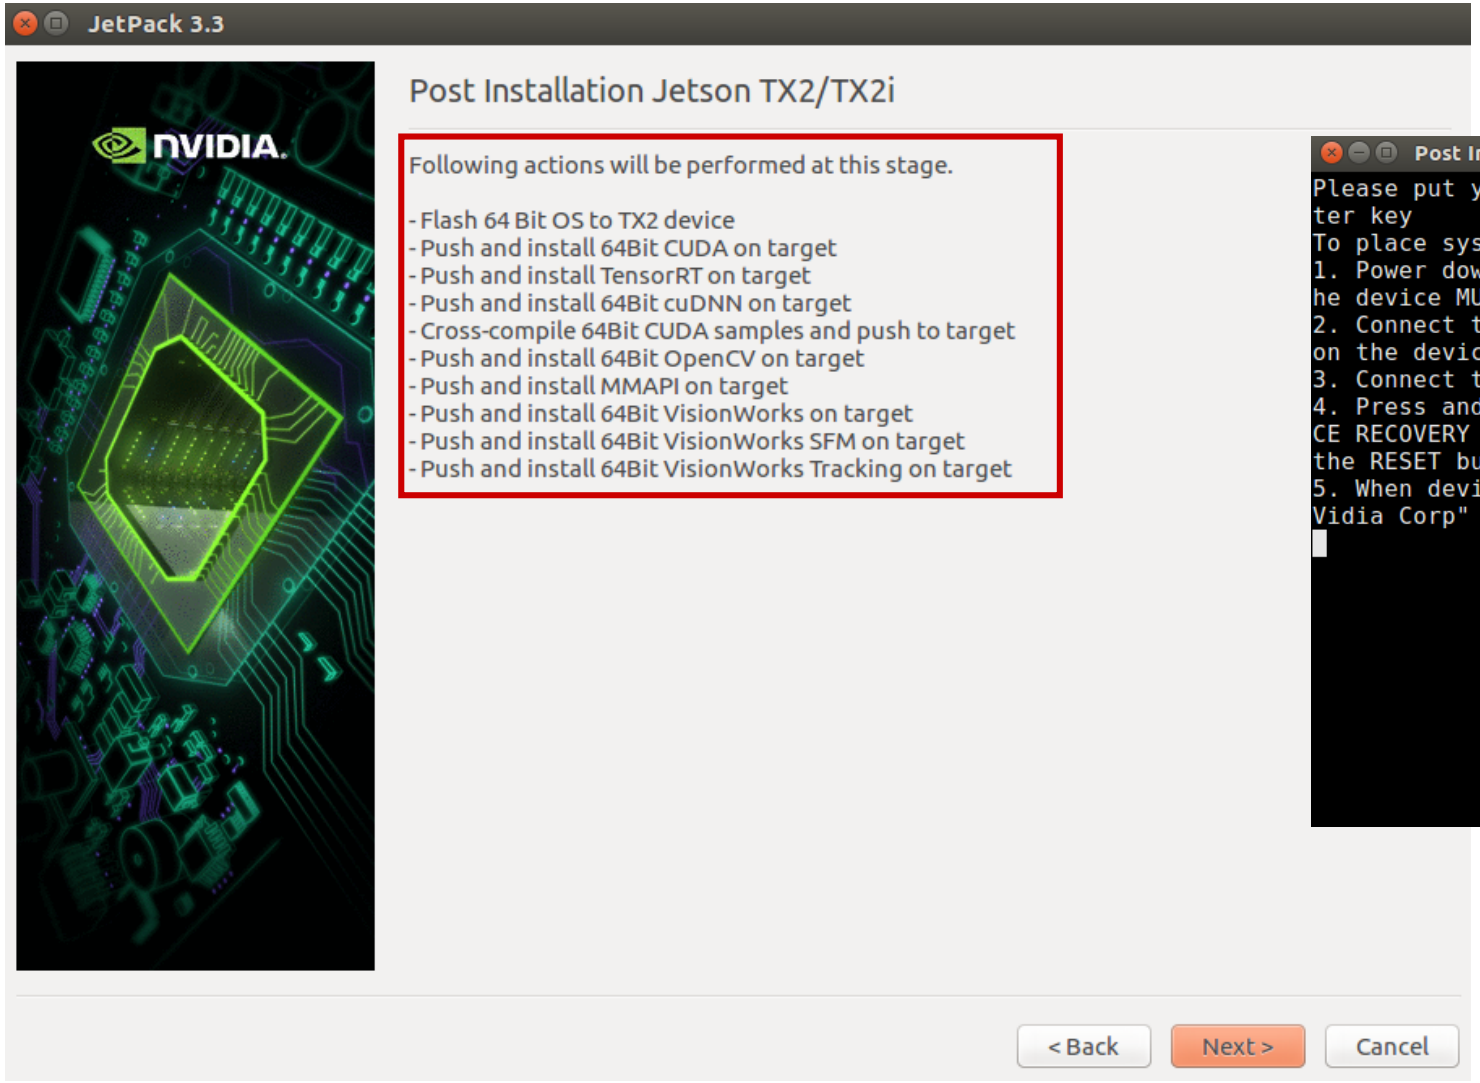

USB Recovery Mode

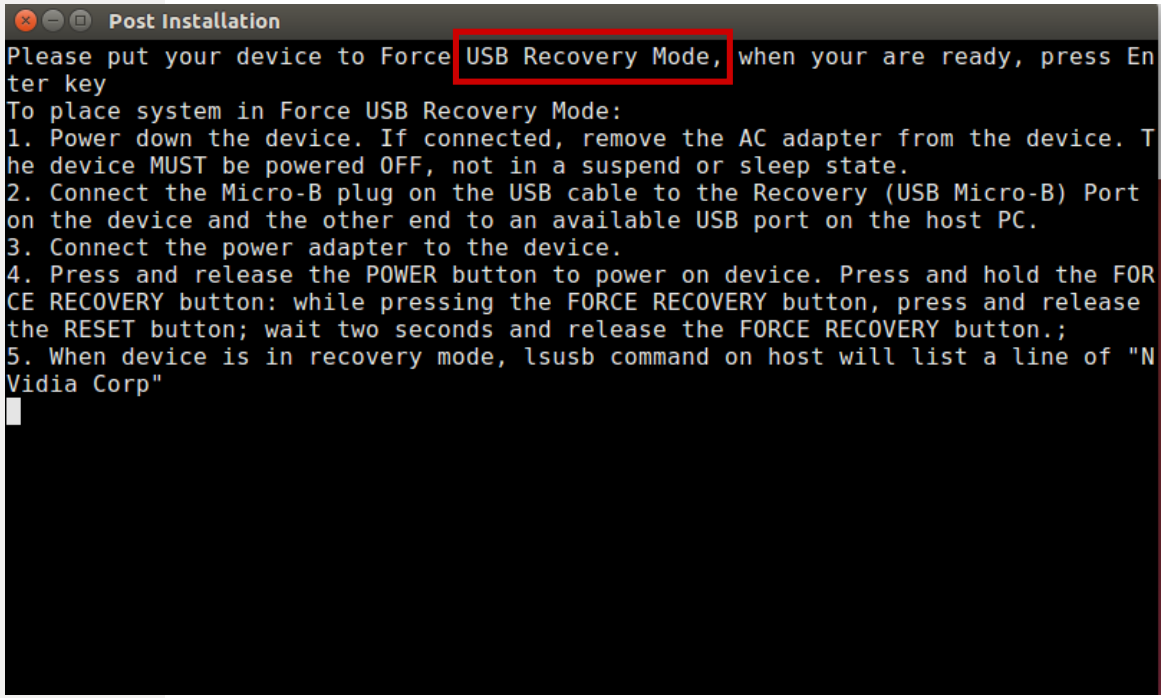

01 JetPack Installation

lsusb

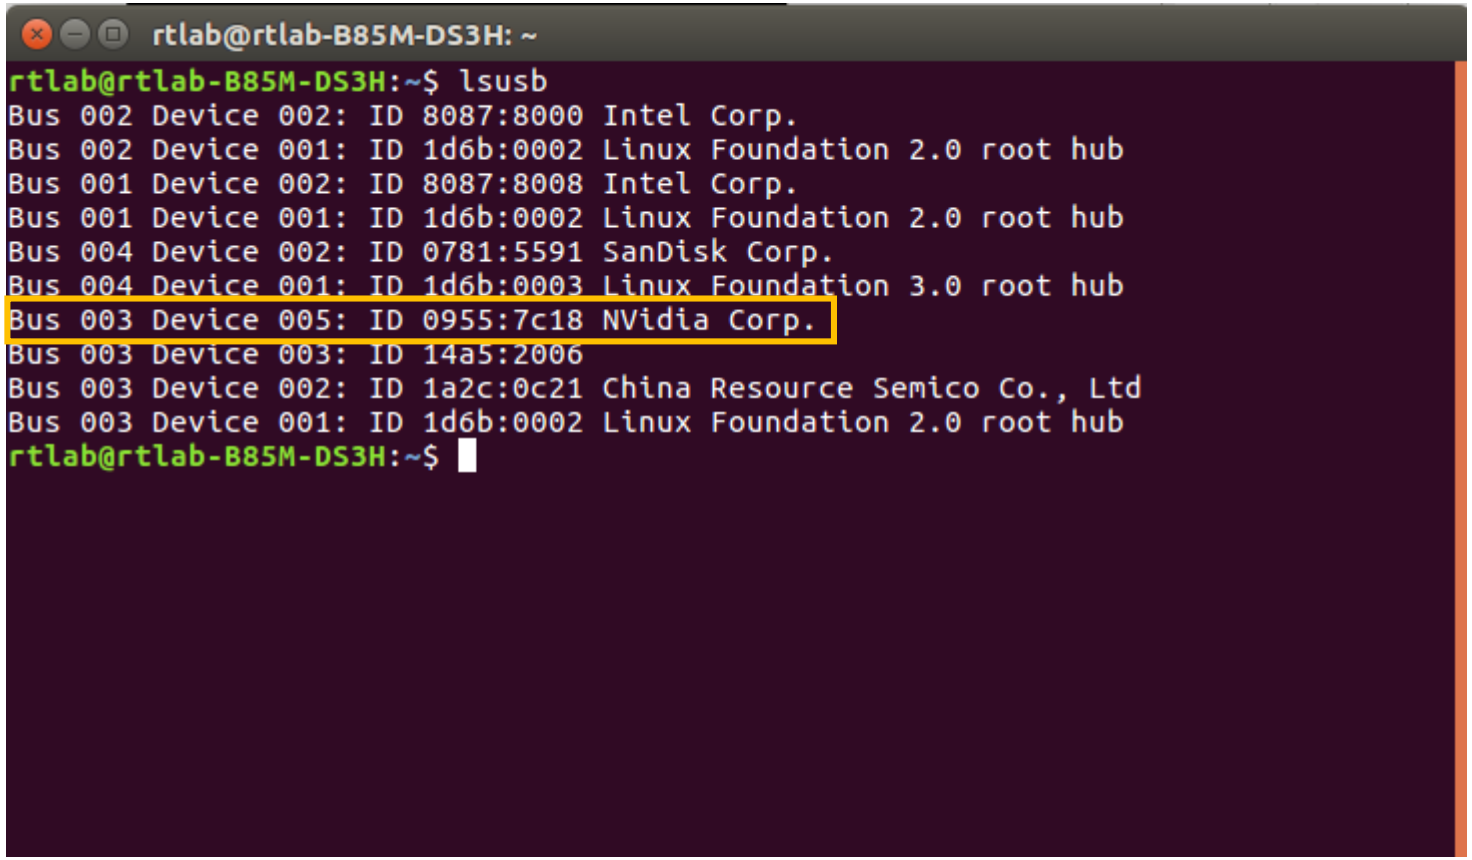

01 JetPack Installation

Building root and flashing.. ~20 minutes

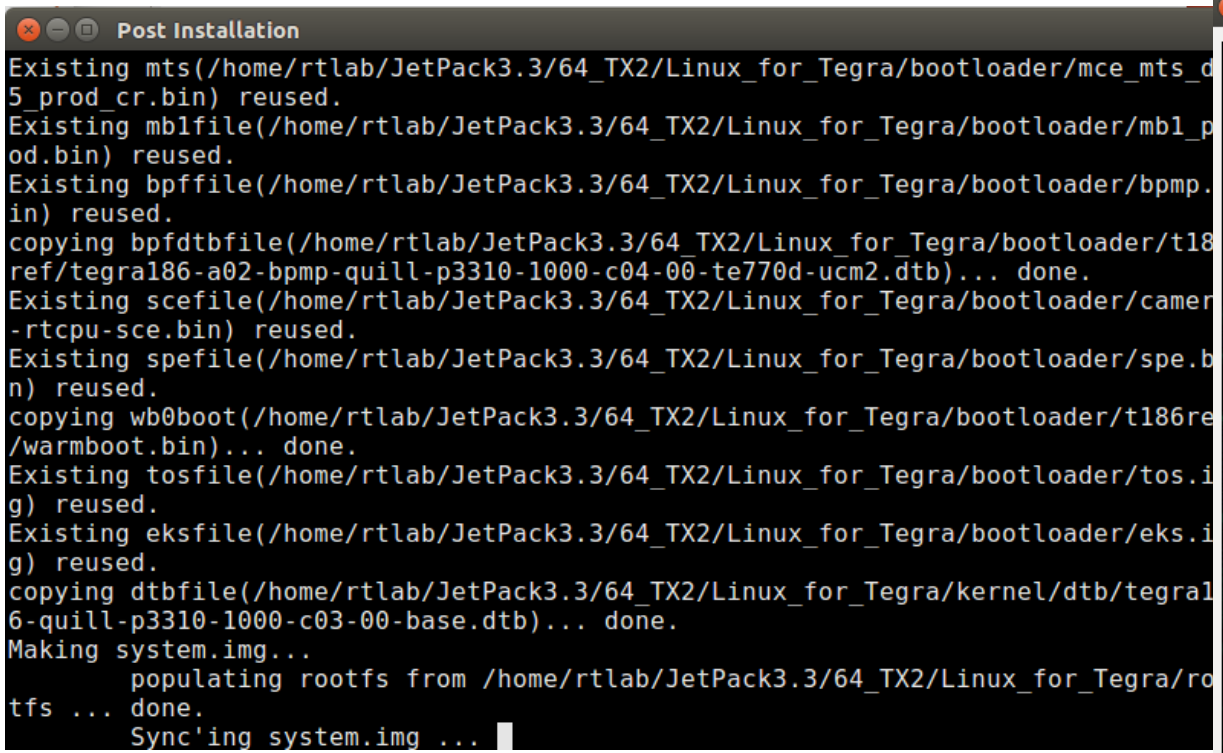

Completion of Target Installation

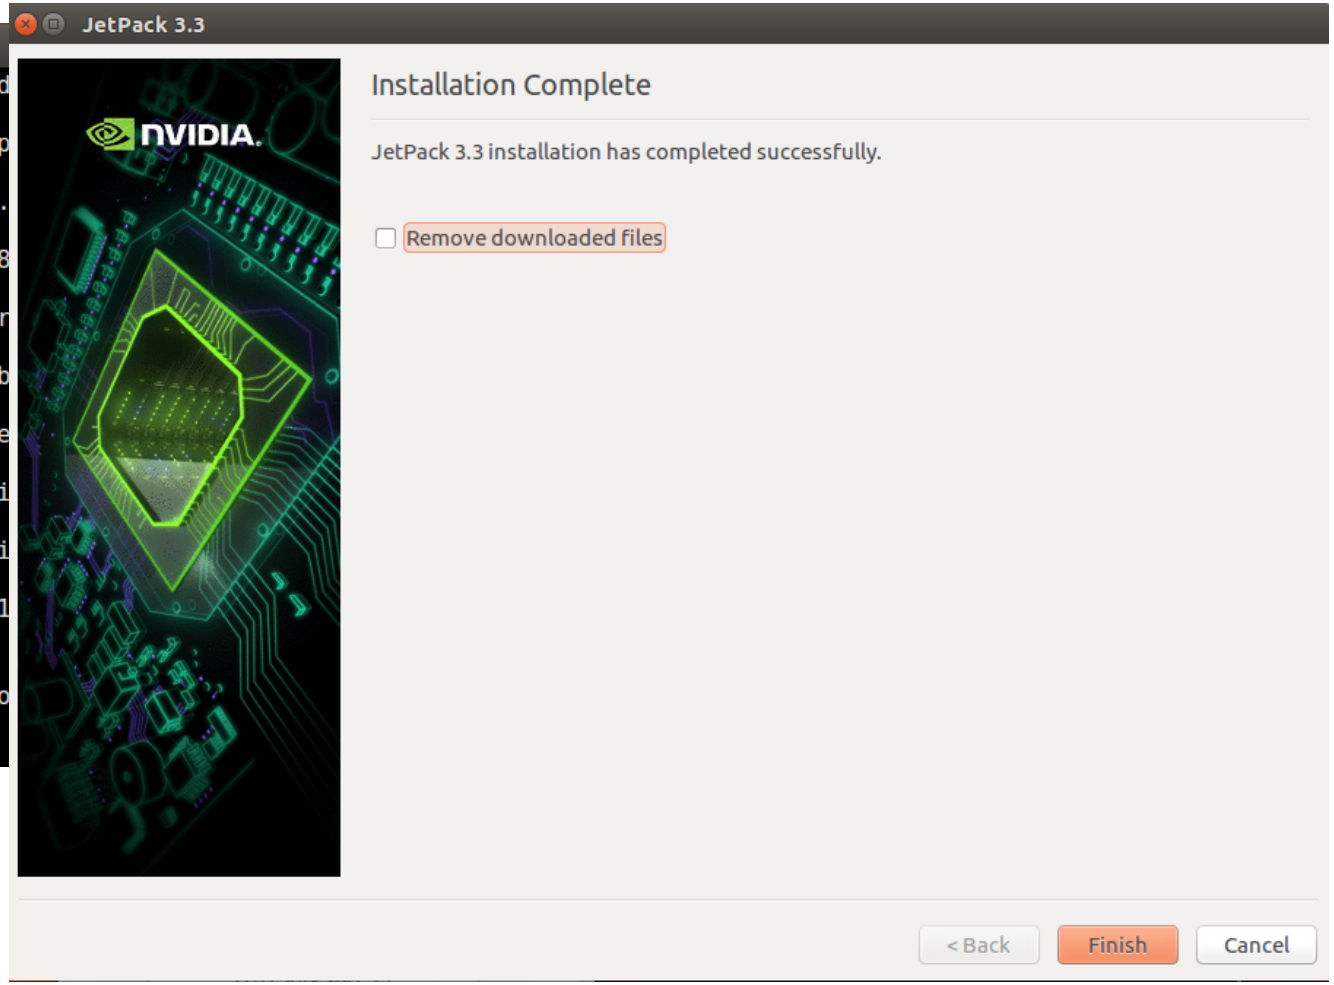

01 JetPack Installation

file:///home/nvidia/L4T/nvl4t_docs/index.html

NVIDIA

L4T Multimedia API Reference

31.0.1 Release

Welcome | Important Terms | API Modules | Sample Applications | Tools | Data Structures | Files | More | Search

- L4T Multimedia API Reference
  - Important Terms
  - API Modules
  - Sample Applications
    - Building and Running
    - Cross-Compiler Support
    - 02_video_dec_cuda
    - 03_video_cuda_enc
    - 04_video_dec_trt
    - 05_jpeg_encode
    - 06_jpeg_decode
    - 07_video_convert
    - 08_video_dec_drm
    - 09_camera_jpeg_capture
    - 10_camera_recording
    - 12_camera_v4l2_cuda
    - 13_multi_camera
    - Backend
    - Frontend
    - v4l2cuda (capture-cuda)
  - Tools
  - Data Structures
  - Files
  - More

Backend

- Overview
- Building and Running
- Flow
- Command Line Options
- Key Structure and Classes

Overview

This application implements a typical appliance performing intelligent video analytics. Application areas include public safety, smart cities, and autonomous machines. This example demonstrates four (4) concurrent video streams going through a decoding process using the on-chip decoders, video scaling using on chip scalar, and GPU compute. For simplicity of demonstration, only one of the channels uses NVIDIA® TensorRT™ to perform object identification and generate bounding box around the identified object. This sample also uses video converter functions for various format conversions. It also uses EGLImage to demonstrate buffer sharing and image display.

In this sample, object detection is limited to identifying cars in video streams of 960 x 540 resolution, running up to 14 FPS. The network is based on GoogleNet. The inference is performed on a frame-by-frame basis and no object tracking is involved. Note that this network is intended to be an example that shows how to use TensorRT to quickly build the compute pipeline. The sample includes trained GoogleNet, which was trained with NVIDIA Deep Learning GPU Training System (DIGITS). The training was done with roughly 3000 frames taken from 5-10 feet elevation. Varying levels of detection accuracy are expected based on the video samples fed in. Given that this sample is locked to perform at half-HD resolutions under 10 FPS, video feeds with higher FPS for inference will show stuttering during playback.

This sample does not require a camera.

Building and Running

Prerequisites

- You have followed Steps 1-3 in Building and Running.
- You have installed:
  - CUDA Toolkit
  - OpenCV
- Optionally, you have installed NVIDIA® TensorRT™ (previously known as GPU Inference Engine (GIE))

To build

1. If you want to run the sample without TensorRT, set the following in the Makefile:

```
ENABLETRT := 0
```

   By default, TensorRT is enabled.
2. Enter:

```
$ cd backend
$ make
```

To run

- Enter:

```
$ ./backend 1 ../../data/Video/sample_outdoor_car_1080p_10fps.h264 H264 \
   --trt-deployfile ../../data/Model/GoogleNet_one_class/GoogleNet_modified_oneClass_halfHD.prototxt \
```

01 JetPack Installation

Sample Application Selected

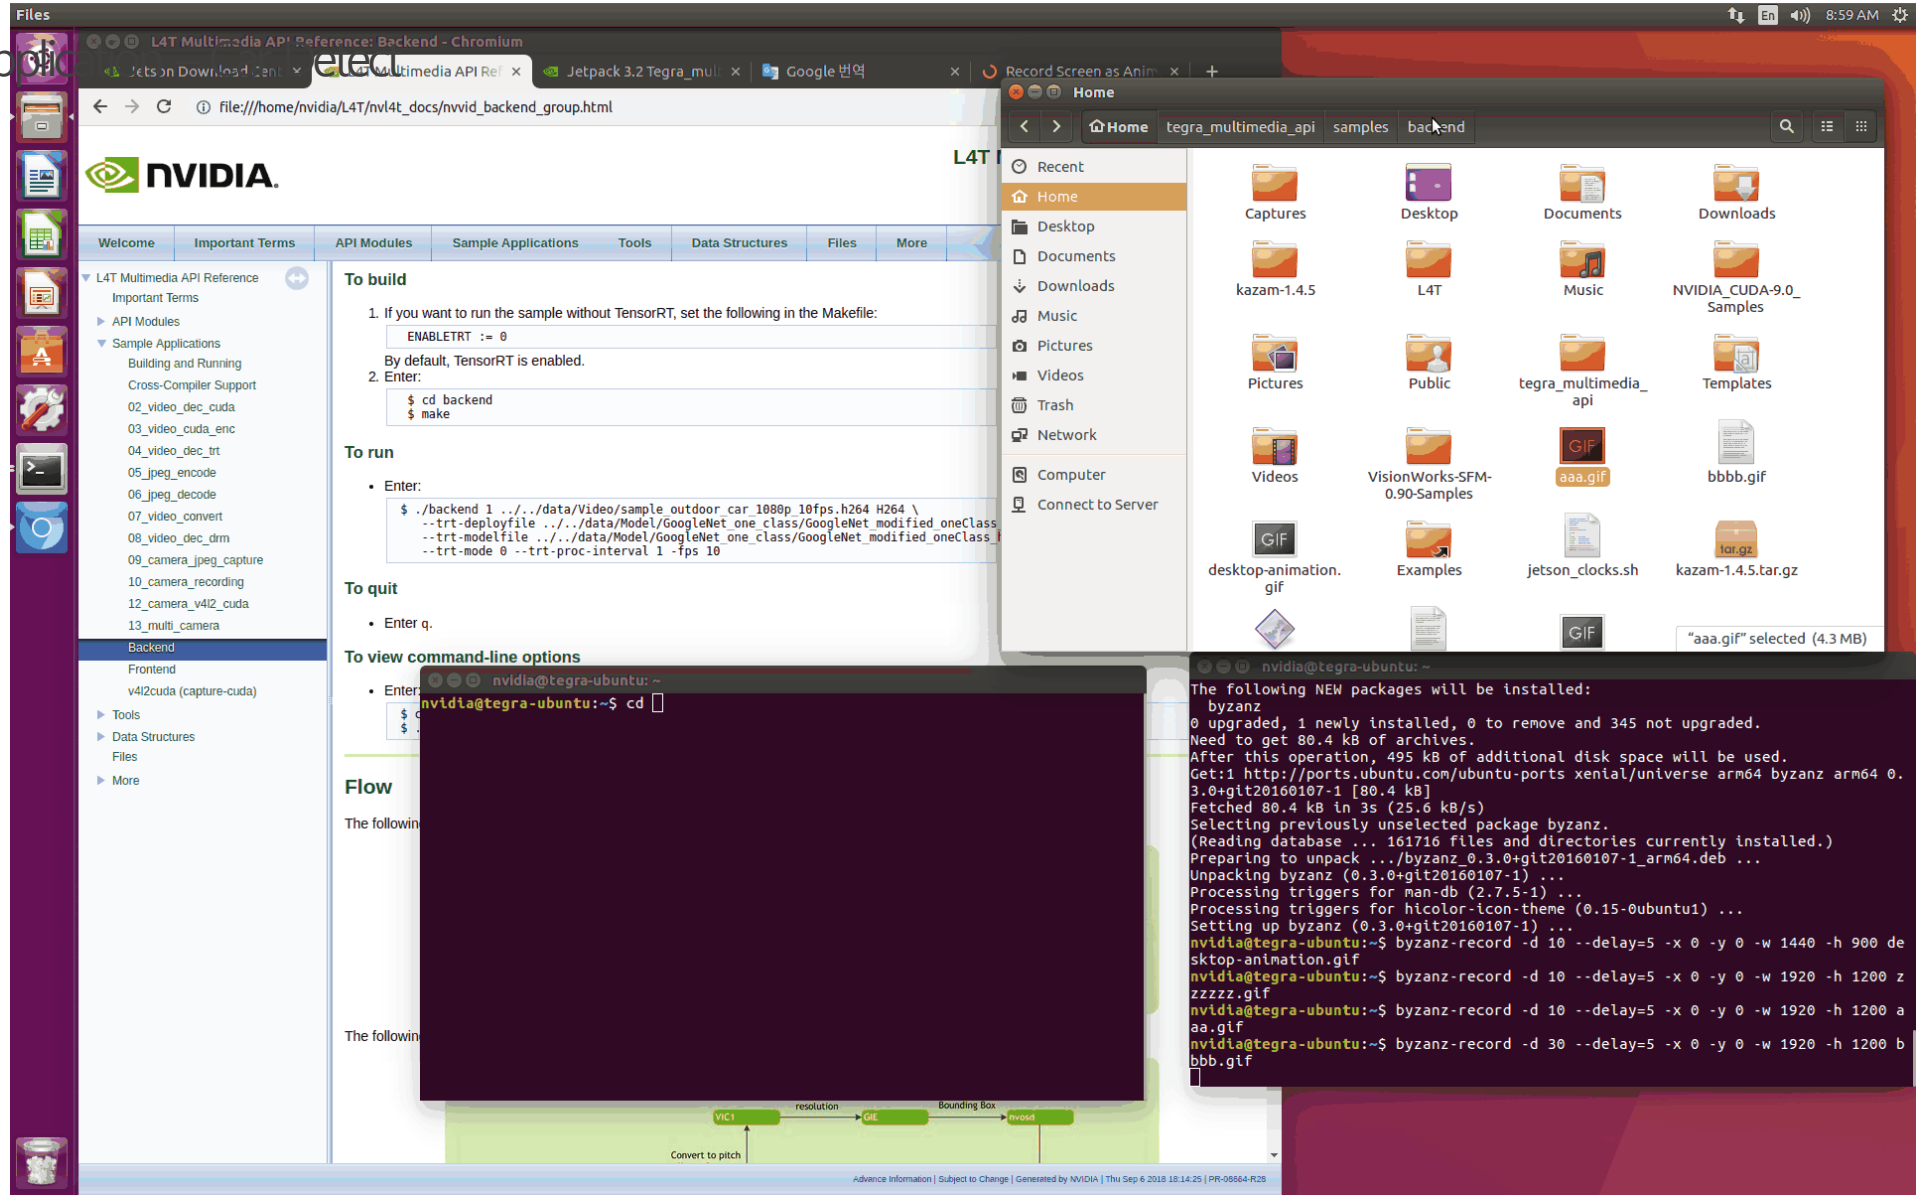

02 Scikit-learn Installation

NVidia-Jetson-TX2-Install-Guide-Lines

https://github.com/Naurislv/NVidia-Jetson-TX2-Install-Guide-Lines

README.md

Preparing Machine Learning/Computer Vision environment for NVidia Jetson TX2

Notes. While installing Ubuntu 16.04.2 LTS on NVidia Jetson TX2 with Tensorflow, Python3.5 and related Computer Vision libraries.

Flashing NVidia Jetson TX2

It is recommended that after unboxing NVidia Jetson TX2 you install fresh OS. NVidia make it very easy by providing JetPack (of-course there are other ways to flash TX2, but this is recommended).

Follow instructions to install latest Ubuntu LTS (for Tegra) on NVidia Jetson TX2 (also aplicable for TX1) using JetPack. While writing this piece JetPack can be installed only on Ubuntu 16.04. There are issues with flashing TX2 with JetPack from VirtualBox and by booting from USB flash drive with Ubuntu 16.04 so I'd recommend to use complete Ubuntu 16.04 installation as Host for JetPack.

You should be able to login using SSH without password but for some operations you may need root priviledges. Default credentials :

username: nvidia ; **password**: nvidia

```
ssh nvidia@ip.address
```

Run fan:

```
sudo ~/jetson_clocks.sh
```

OS details:

02 Scikit-learn Installation

Upgrade Ubuntu

```
sudo apt-get update && sudo apt-get upgrade
```

```
sudo apt-get install python3-numpy swig python3-dev python3-pip python3-wheel -y
```

- Upgrade pip : `sudo pip3 install --upgrade pip`
- Other dependecies (replace 3.5 with the python version you are using). Order is important.:

i. `sudo pip3.5 install --upgrade azure.common` (If you will not use MS Azure ignore this)
ii. `sudo pip3.5 install --upgrade azure.servicebus` (If you will not use MS Azure ignore this)
iii. `sudo pip3.5 install --upgrade azure.storage` (If you will not use MS Azure ignore this)
iv. `sudo pip3.5 install --upgrade setuptools ez_setup`
v. `sudo pip3.5 install --upgrade matplotlib`
vi. `sudo pip3.5 install --upgrade configparser`
vii. `sudo pip3.5 install --upgrade numpy`
viii. `sudo apt-get install libblas-dev liblapack-dev libatlas-base-dev`
ix. `sudo pip3.5 install --upgrade scipy`
x. `sudo pip3.5 install --upgrade pandas`
xi. `sudo pip3.5 install --upgrade sklearn`

Building on Linux

It is easiest to use your system package manager to install the dependencies.

If you are on Debian/Ubuntu, you can get all the dependencies required to build Matplotlib with:

```
sudo apt-get build-dep python-matplotlib
```

02 Scikit-learn Installation

chap3 example code

localhost:8888/notebooks/Untitled.ipynb - Chromium

localhost:8888/notebooks/Untitled.ipynb

jupyter Untitled Last Checkpoint: 42 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted | Python 3

Run Code

In [18]:

```
def plot_digit(data):
    image = data.reshape(28, 28)
    plt.imshow(image, cmap = matplotlib.cm.binary,
               interpolation="nearest")
    plt.axis("off")
```

In [19]:

```
# 숫자 그림을 위한 추가 함수
def plot_digits(instances, images_per_row=10, **options):
    size = 28
    images_per_row = min(len(instances), images_per_row)
    images = [instance.reshape(size,size) for instance in instances]
    n_rows = (len(instances) - 1) // images_per_row + 1
    row_images = []
    n_empty = n_rows * images_per_row - len(instances)
    images.append(np.zeros((size, size * n_empty)))
    for row in range(n_rows):
        rimages = images[row * images_per_row : (row + 1) * images_per_row]
        row_images.append(np.concatenate(rimages, axis=1))
    image = np.concatenate(row_images, axis=0)
    plt.imshow(image, cmap = matplotlib.cm.binary, **options)
    plt.axis("off")
```

In [26]:

```
plt.figure(figsize=(9,9))
example_images = np.r_[X[:12000:600], X[13000:30600:600], X[30600:60000:590]]
plot_digits(example_images, images_per_row=10)

plt.show()
```

In [21]:

03 NVIDIA Two Days to a Demo

https://developer.nvidia.com/embedded/twodaystoademo

Two Days to a Demo

Two Days to a Demo is our introductory series of deep learning tutorials for deploying AI and computer vision to the field with NVIDIA Jetson AGX Xavier, Jetson TX1 and Jetson TX2. This tutorial takes roughly two days to complete from start to finish, enabling you to configure and train your own neural networks. It includes all of the necessary source code, datasets, and documentation to get you started. Dive into deep learning today with Two Days to a Demo.

Deploying Deep Learning

Ready to dive into deep learning? It only takes two days. We'll provide you with all the tools you need, including easy to follow guides, software samples such as TensorRT code, and even pre-trained network models including ImageNet and DetectNet examples. Follow these directions to integrate deep learning into your platform of choice and quickly develop a proof-of-concept design. In this guide, you'll get a stronger background in deep learning, be able to load and run a pre-trained deep neural network on the Jetson AGX Xavier Developer Kit or Jetson TX1/TX2 Developer Kit, and learn how to retrain the network with your own dataset to produce a live demo.

Four Steps to Deep Learning

1. System Setup
2. Image Recognition
3. Object Detection
4. Segmentation

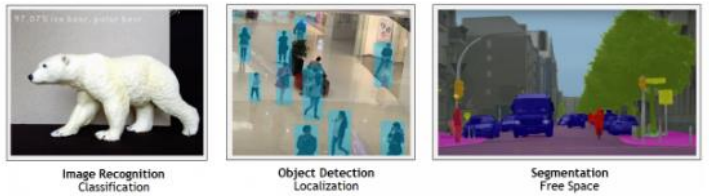

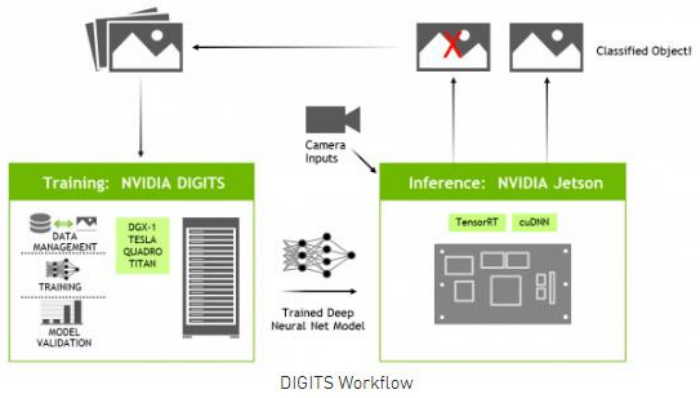

DIGITS Workflow

03 NVIDIA Two Days to a Demo

1. Installing Ubuntu on the Host
2. Running JetPack on the Host
3. Installing the NVIDIA driver (v390.87)
4. Installing Docker

```
docker run --runtime=nvidia --rm nvcr.io/nvidia/cuda:9.0-cudnn7-devel-ubuntu16.04 nvidia-smi
```

Table 1. CUDA Toolkit and Compatible Driver Versions

CUDA Toolkit	Linux x86_64 Driver Version	Windows x86_64 Driver Version
CUDA 10.0.130	>= 410.48	>= 411.31
CUDA 9.2 (9.2.148 Update 1)	>= 396.37	>= 398.26
CUDA 9.2 (9.2.88)	>= 396.26	>= 397.44
CUDA 9.1 (9.1.85)	>= 390.46	>= 391.29
CUDA 9.0 (9.0.76)	>= 384.81	>= 385.54
CUDA 8.0 (8.0.61 GA2)	>= 375.26	>= 376.51
CUDA 8.0 (8.0.44)	>= 367.48	>= 369.30
CUDA 7.5 (7.5.16)	>= 352.31	>= 353.66
CUDA 7.0 (7.0.28)	>= 346.46	>= 347.62

Thank You!

01 Jetpack Install

자세한 내용을 입력해 주세요

DAILY-C

데일리C 레몬워터

소중한 내 몸을 위한 특별한 물

레몬의 상큼함이 듬뿍!

물보다 맛있는 음료 데일리C 레몬워터

비타민C가 함유된 리프레시 비타민 음료

240ml캔, 500ml 페트

DAILY-C

데일리C 레몬1000C+

레몬 과즙이 들어있는 리얼 레몬 드링크

하루 한 병으로 비타민C를 상큼하게 충전하는

리얼 레몬 드링크

레몬 32개 분량에 해당하는 비타민C 1,0000mg

을

가볍게 충전해 보세요

140ml 병

그래서 우리 발표

아이디어를

이렇게 글씨로

포인트 살려서 예쁘게

적어 주세요